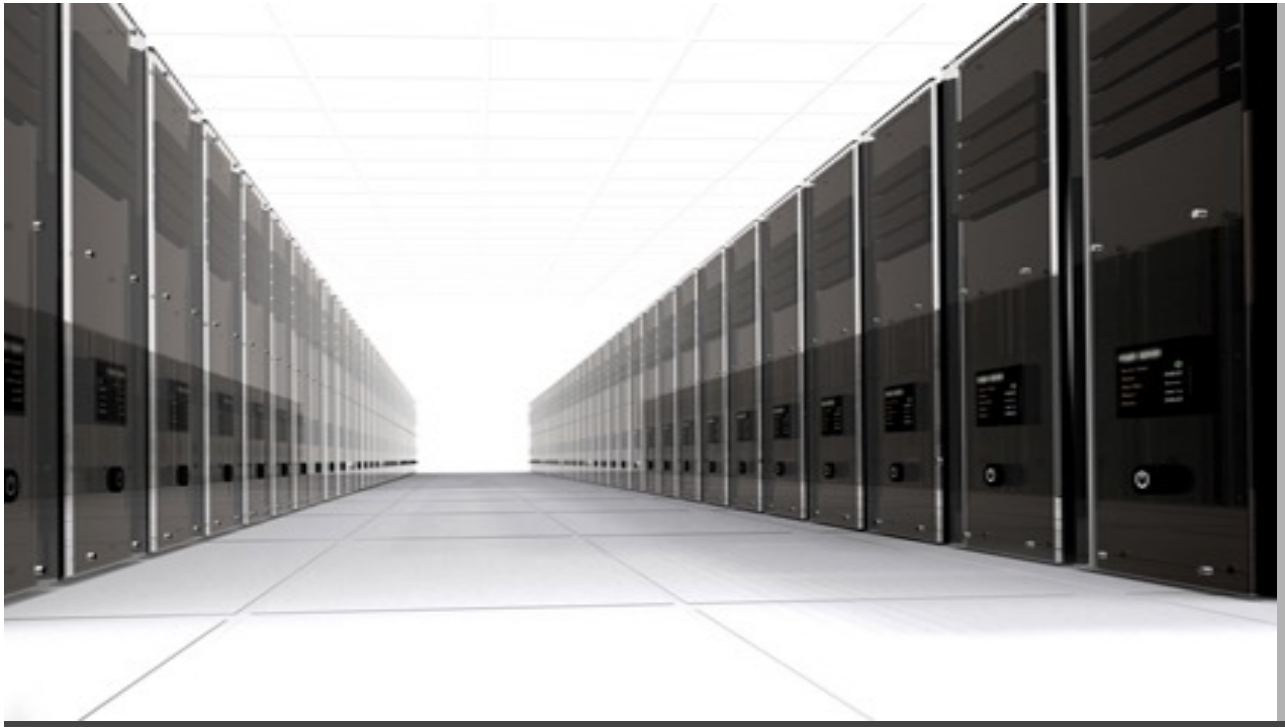

STORAGE (II): BASIC CONCEPTS

David López

v 2.6.1

Updated fall 2024

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

Magnetic vs. Optical vs. Solid State

Three basic storage technology:

- Magnetic
  - Tapes (1952-Today)
  - Hard Disk (1956-Today)
- Optical
  - Optical Disc Archive (2013 – Today)
- Solid State
  - Solid State Discs – SSD (2006 – Today)

Hard Disk situation

Hard disks are "living dinosaurs"

- According to Moore's law, the density of microelectronics doubles every 18 months
- In hard disks, this only applies to:
  - Process speed of the controller (which never was much of a problem anyway)
  - Increased speed of read/write operations because more data is packed onto each track
  - Increased capacity of the disk (that means more accesses per second)
- The problem is that it does not affect nor to the rotational speed neither to the actuators moving speed
  - And several actuators on the same rack does not work due to the high density and dilatation

BIG PROBLEM: HDD can store gigantic amounts of data, but the transactions per second are tied to the mechanical internals

Hard disk problems

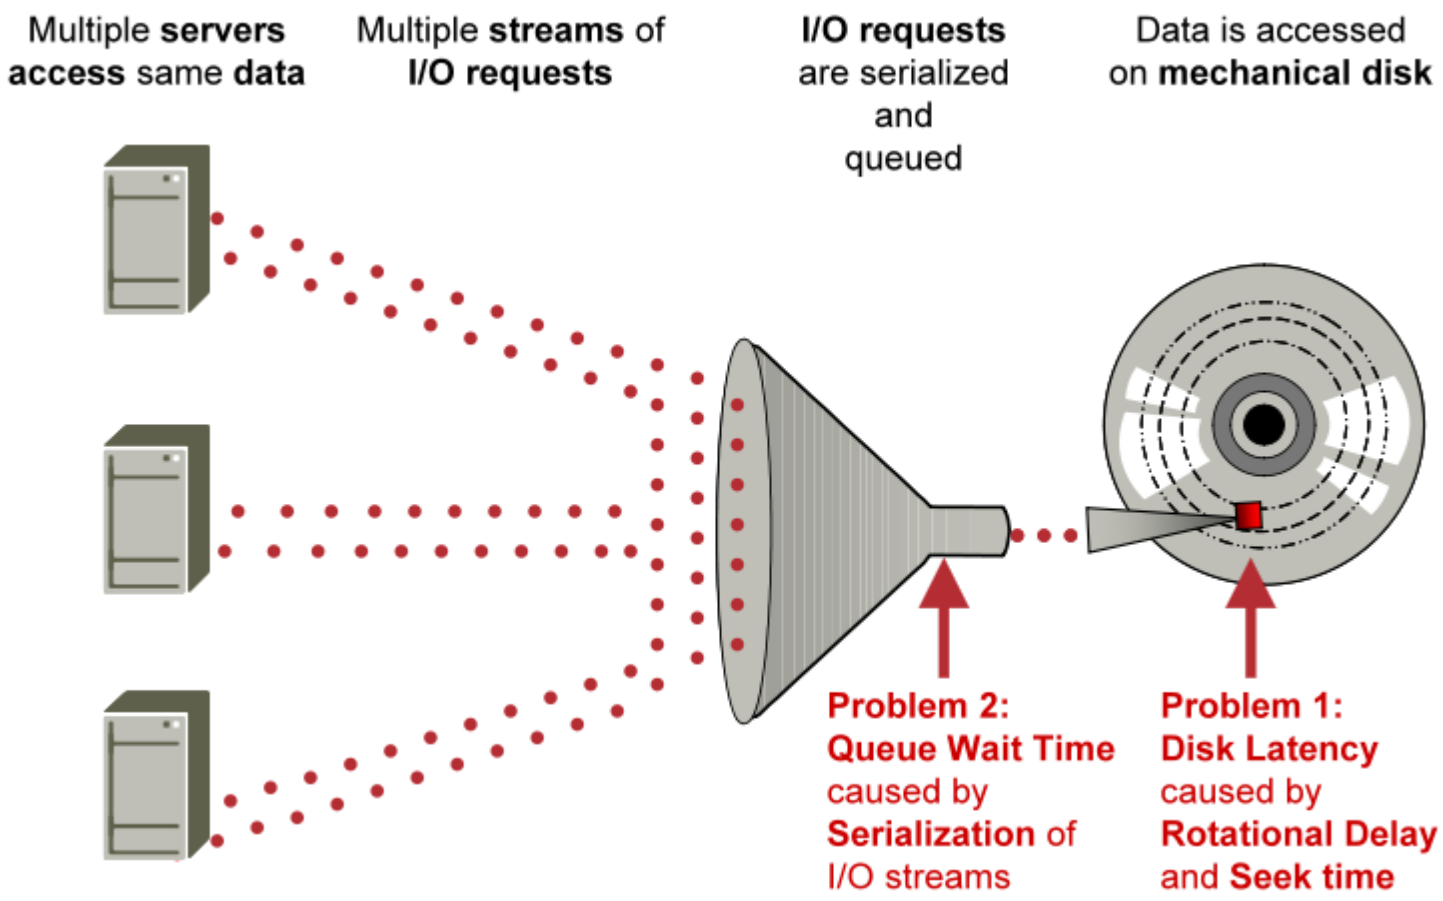

http://www.violin-memory.com/assets/Violin-WP-Disk-Storage-Shortfall.pdf?d=1

Disk latency + queue wait time

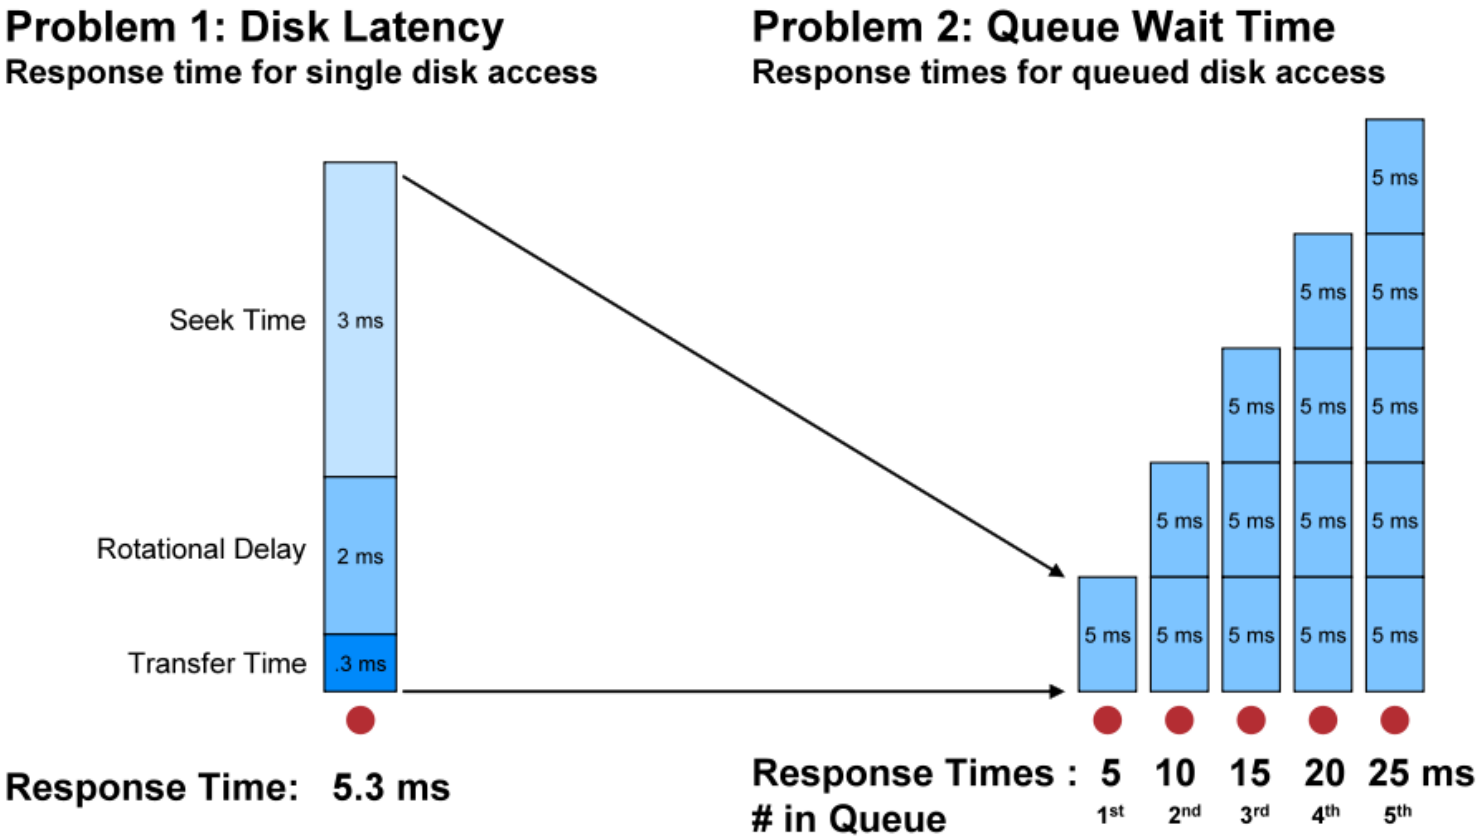

http://www.violin-memory.com/assets/Violin-WP-Disk-Storage-Shortfall.pdf?d=1

Storage triangle

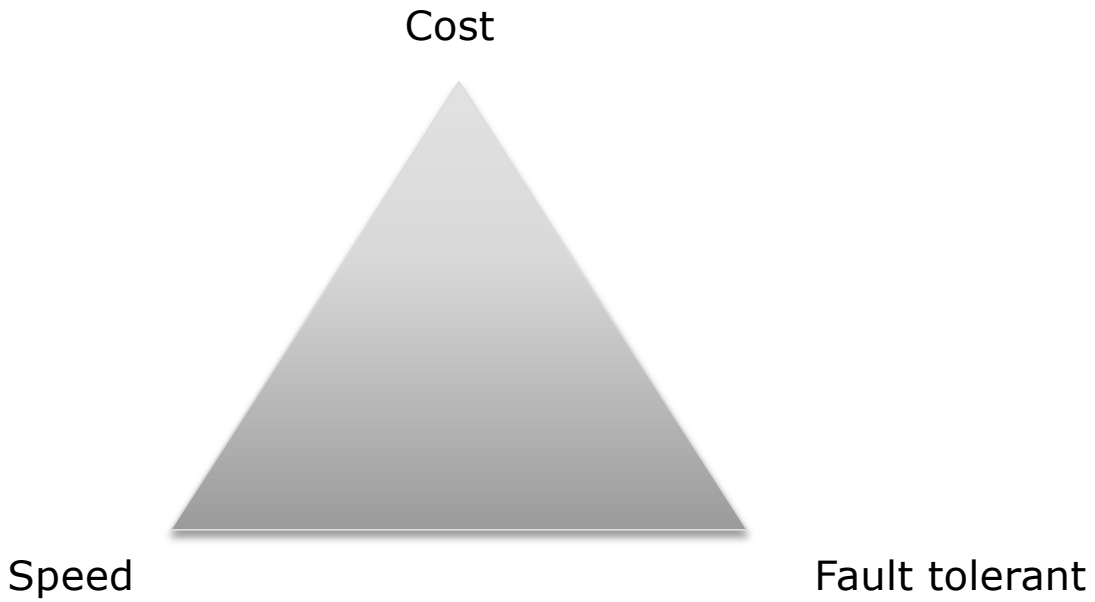

LUNs and JBOD

- Divided in **LUN**s (**L**ogical **UN**its)
  - For the host computer, there are not differences between LUNs and physical disks
- Easy to work for the host computer
  - Partitions or (more often) aggregation
  - Saw as an unique disk for backup
- Example a **JBOD** (**J**ust a **B**unch **O**f **D**isks)
  - Example: three 2TB disks
  - Build a 6TB LUN
  - You can have disks of different size (not like RAID)
  - One block following the next on the same disk (not like RAID 0) "Concatenation or SPAN, not stripped"

JBOD

	Space Efficiency	Fault tolerance	Read Performance	Write Performance
JBOD	1	0	1	1

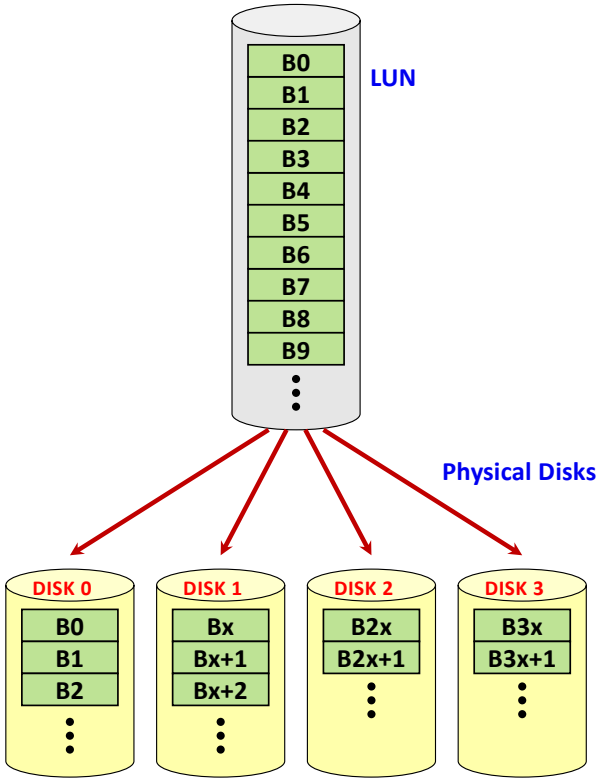

Image by Agustín Fernández (AC)

Storage triangle

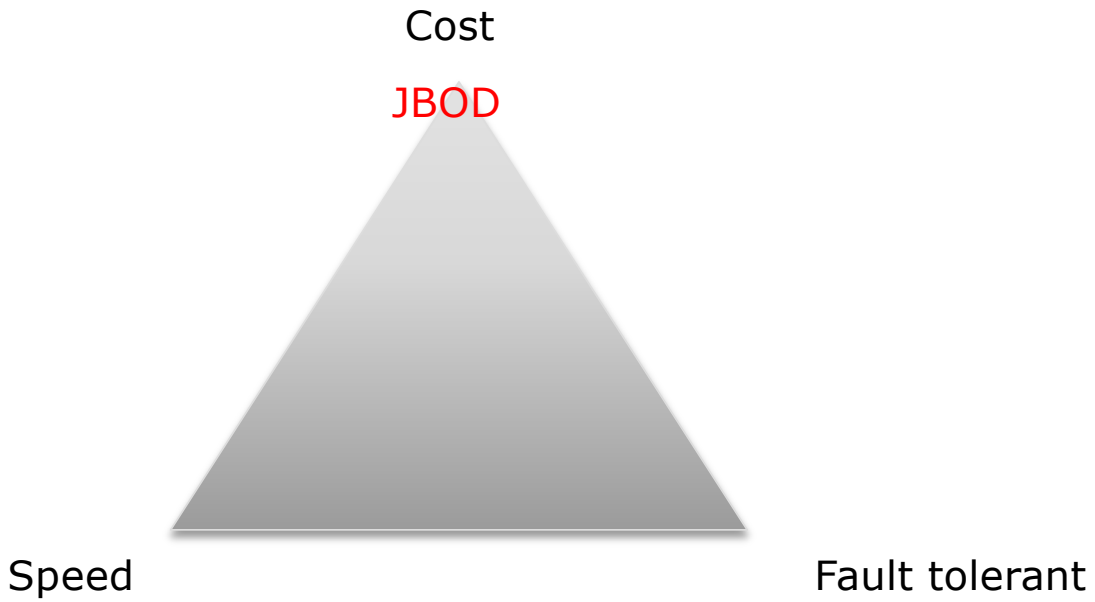

Avoiding errors: RAID

- RAID offers redundancy, BUT ALSO SPEED (at a certain cost)
- Let's calculate # of parallel R/W in
  - RAID 0
  - RAID 1
  - RAID 5
  - RAID 6
  - RAID 10, 01
  - RAID 51, 15
- Important question: WHAT ABOUT THE STRIPE SIZE?
  - 4KB-128KB?
  - In the activities we will consider 4KB but it is an interesting question

RAID 0 (stripping) & RAID 1 (mirroring)

	Space Efficiency	Fault tolerance	Read Performance	Write Performance
RAID 0	1	0	n to 1	n to 1
RAID 1	1/n	n-1	n (real)	1

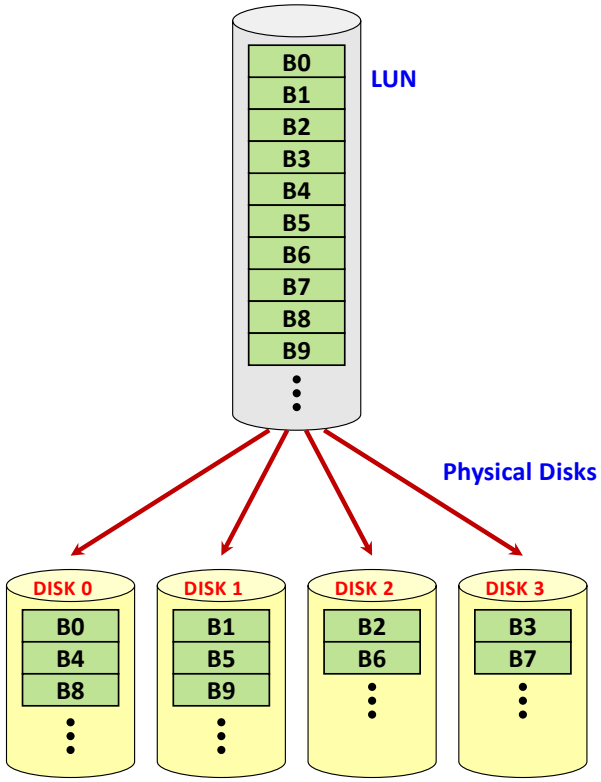

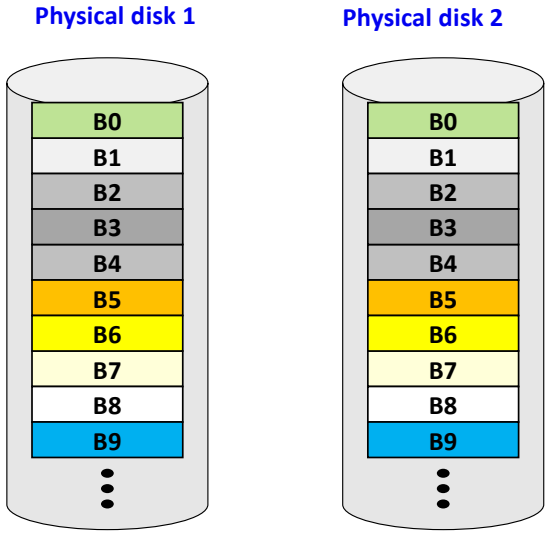

Image by Agustín Fernández (AC)

Storage triangle

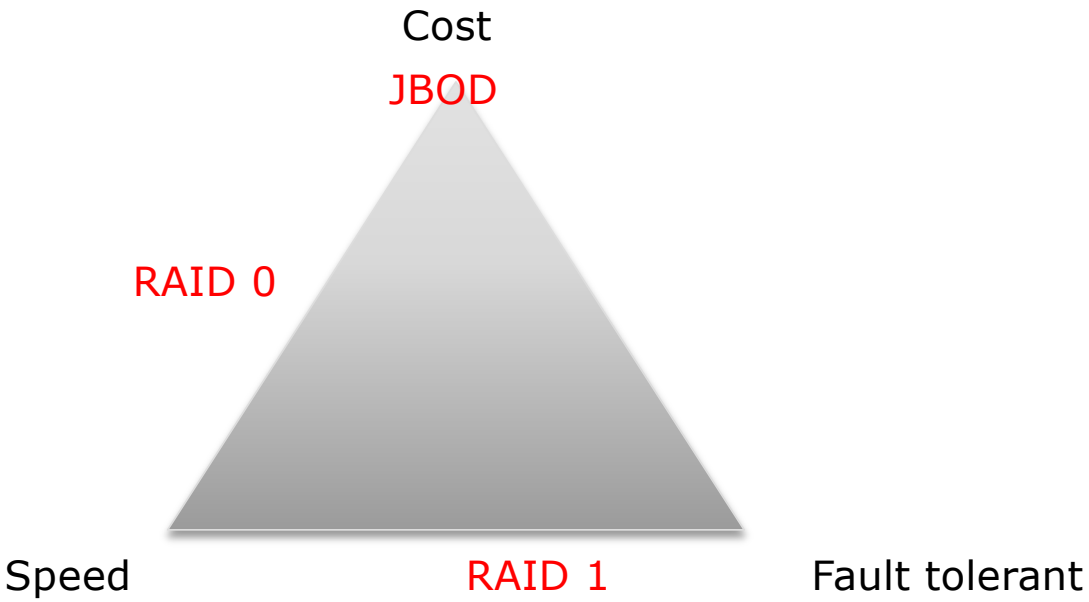

RAID 5: Block-level striping with distributed parity

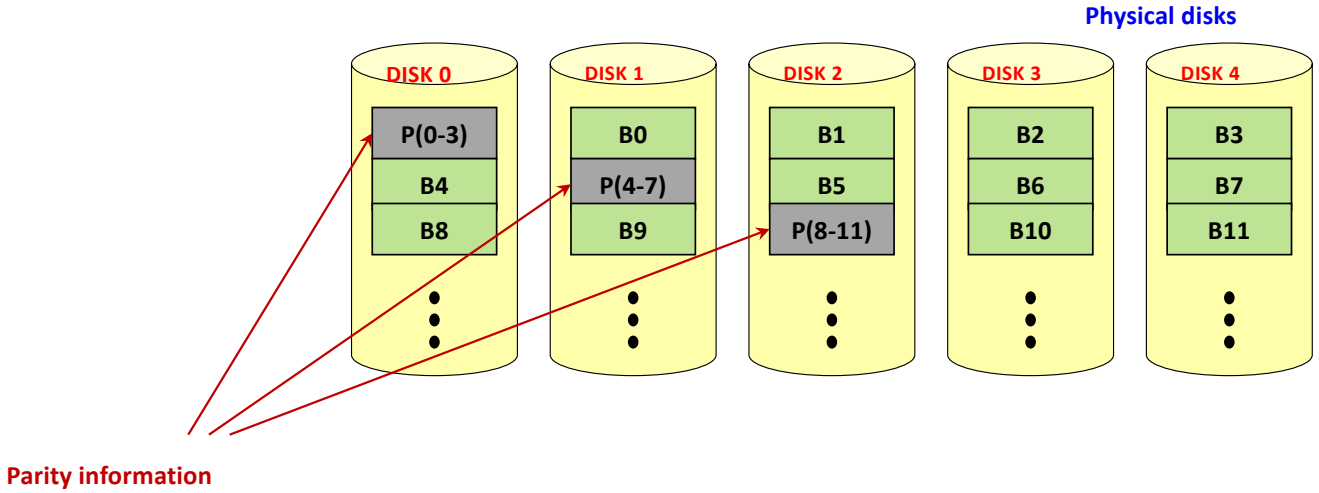

	Space Efficiency	Fault tolerance	Read Performance	Write Performance
RAID 5	n-1	1	n (n/2)	(n-1) (n/2)

Image by Agustín Fernández (AC)

RAID 6: Block-level striping with double distributed parity

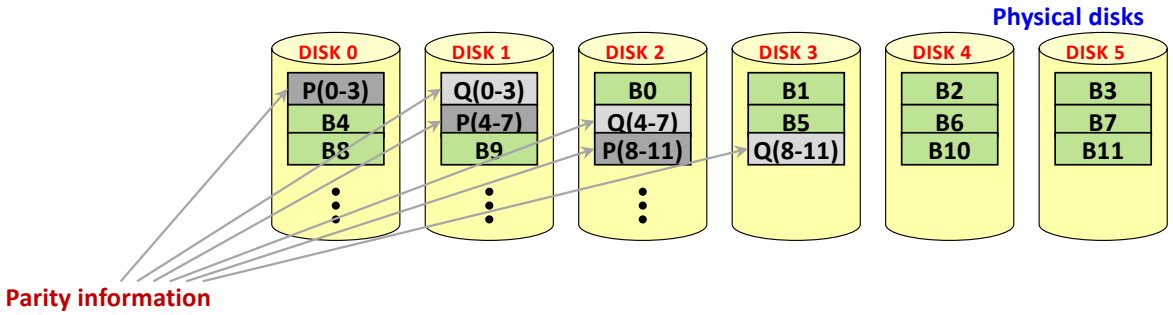

	Space Efficiency	Fault tolerance	Read Performance	Write Performance
RAID 6	n-2	2	n (n/3)	(n-2) (n/3)

Image by Agustín Fernández (AC)

RAID 10 & RAID 01

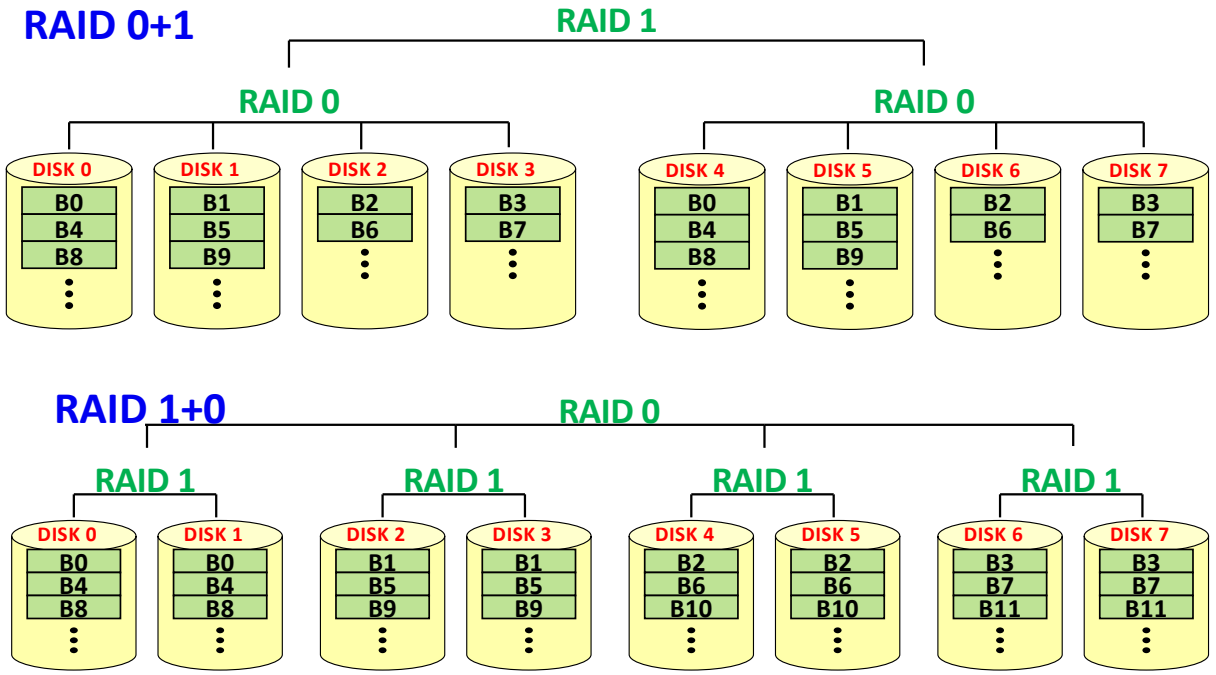

	Space Efficiency	Fault tolerance	Read Performance	Write Performance
RAID 10/01	n/mirrors	n/mirrors	n mirrors	(n/mirrors) 1

Image by Agustín Fernández (AC)

RAID 51 & 15

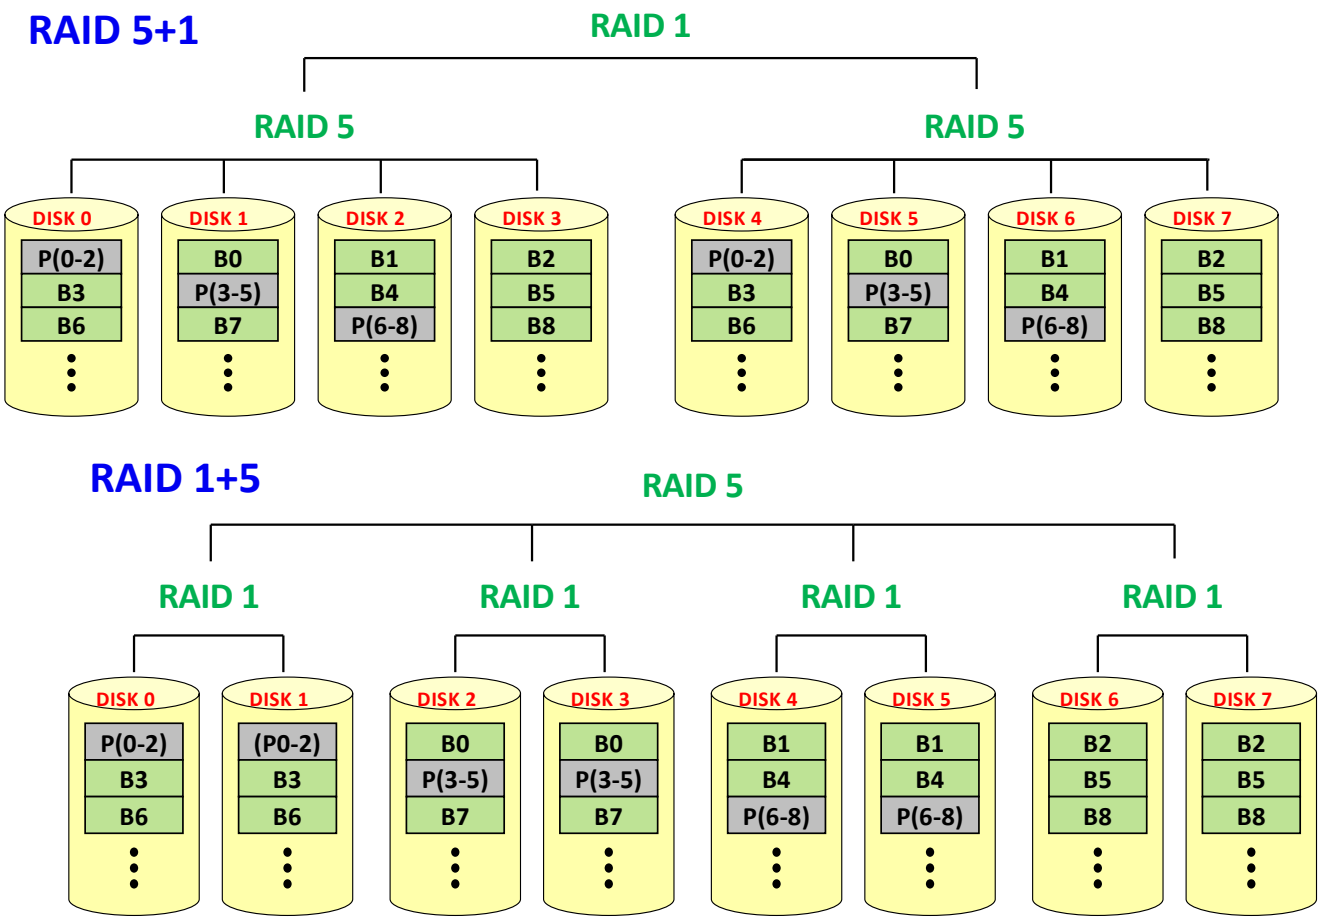

Image by Agustín Fernández (AC)

Storage triangle

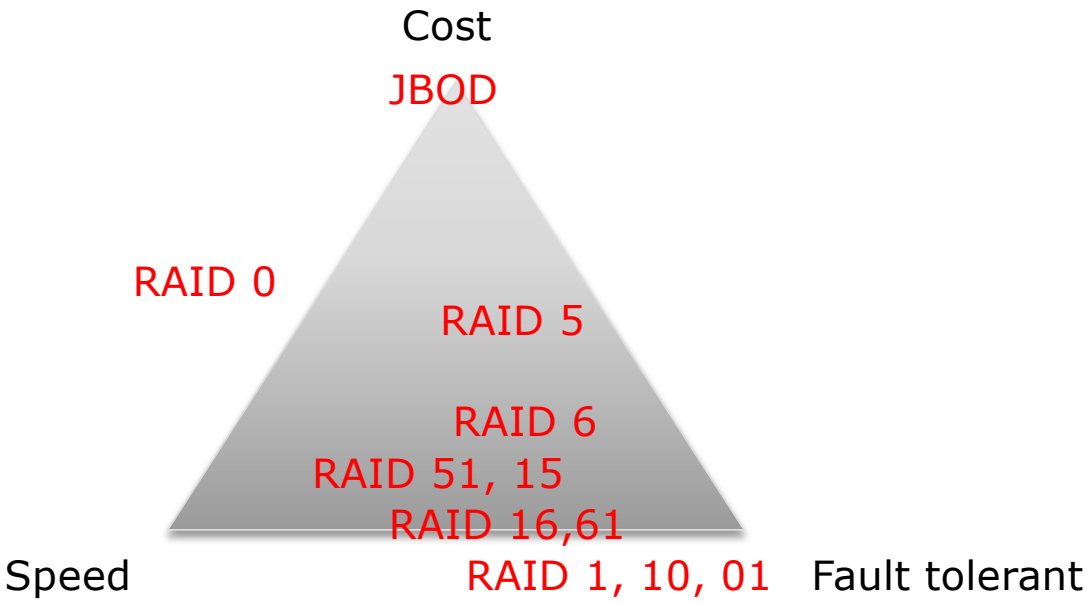

RAID, write penalty & capacity

	RAID 0	RAID 10	RAID 5	RAID 51	RAID 6	RAID 61
Operations per write	1W	2W	2R+2W	(2R+2W) x2	3R+3W	(3R+3W) x2
Write penalty	1	2	4	8	6	12
Capacity	X*C	(X/2)*C	(X-1)*C	((X-1)/2)*C	(X-2)*C	((X-2)/2*C
Minimum number of discs	2	4	3	6	4	8
Required discs (for Y Bytes)	Y/C	2*Y/C	Y/C +1	2*Y/C +1	Y/C +2	2*Y/C +2

Let's assume X discs, homogeneous, each one of capacity C

IOPS (Input / Output Operations Per Second)

- Pronounced *eye-ops*
- Common performance measurement for storage devices
- There are applications to measure it
  - Iometer (Intel)
  - IOzone
  - FIO
- Not easy to define / compare
  - Mix of read / write operations
  - Sequential and random accesses
  - Data block sizes
- Typical values
  - Total IOPS (mix of R/W, Seq/RND)
  - Random read IOPS
  - Random write IOPS
  - Sequential read IOPS
  - Sequential write IOPS
- IOPS * TransferSizeInBytes = MBps

SSD performance

Many IOPS? Solid State Disks can offer the solution!

- In our project
  - HDD IOPS: 640 – 5210
  - SSD IOPS (RD/WR): 90k/10k – 540k / 205k

And the cost? Fa$t di$k$ co$t money!

- In our project
  - HDD cost: 0,029 /G (8 TB=235€) – 0,15€/G (2.4TB=360€)
  - SSD cost: 0,155 €/GB (2TB=310€) – 0,21€/GB (7,68TB=1545€)

SSD & HDD price forecast

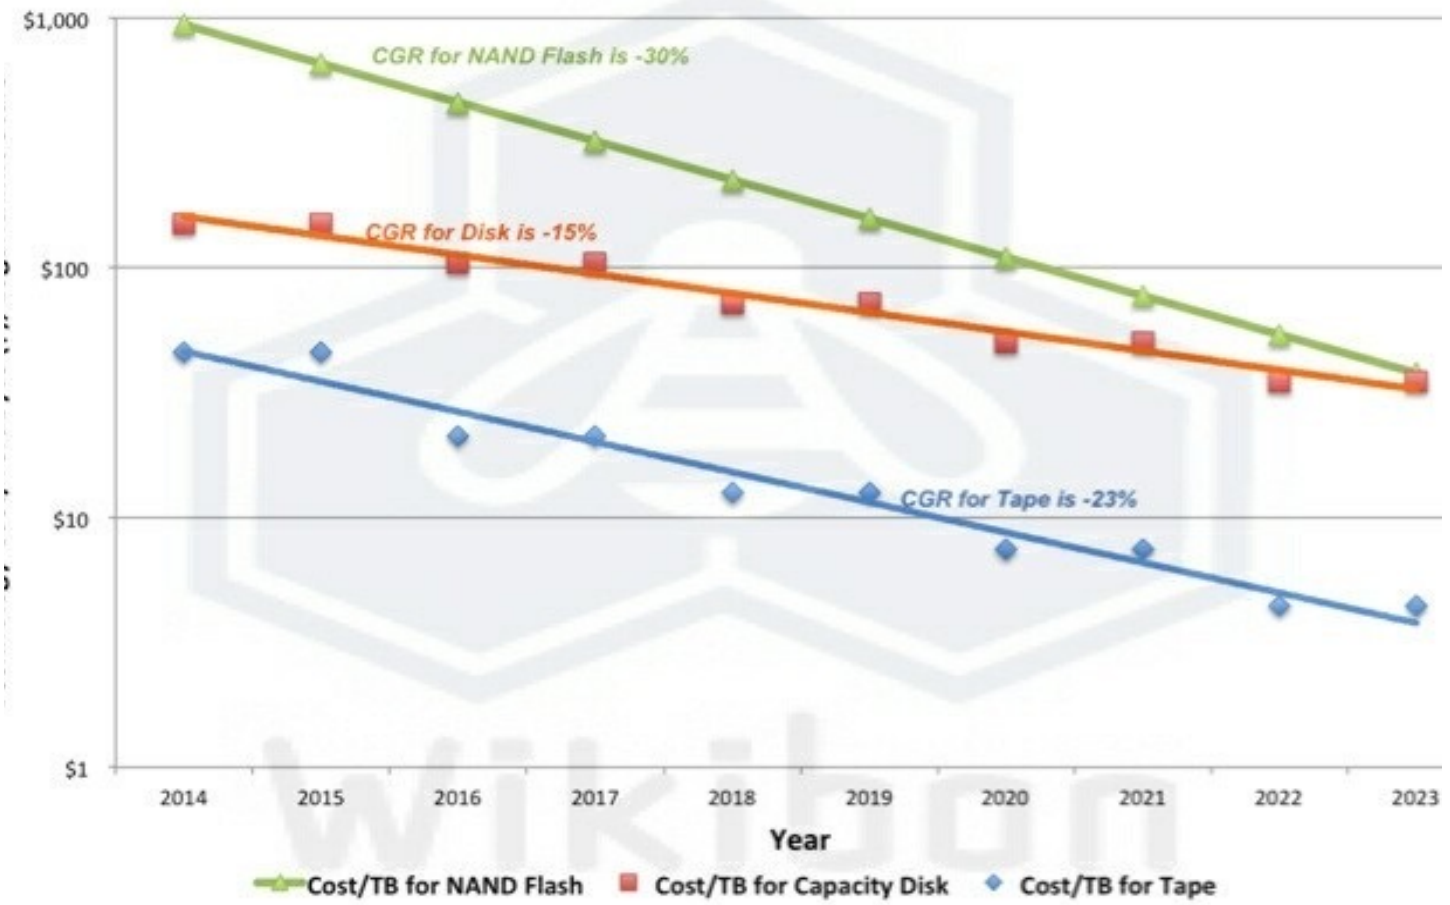

SSD and HDD with similar price?

1. Enterprise vs Consumer
   - Different quality, different costs
2. New technology in magnetic disks
   - Microwave-Assisted Magnetic Recording (MAMR)
     - Using a Spin-Torque Oscilator (STO) to generate microwaves in very small particules
     - 4 Tbpsi (Terabit per square inch)
     - Current perpendicular recording up to 1,31 Tbpsi
   - Heat-Assisted Magnetic Recording (HAMR)
     - Using laser to heat the material
     - 5 Tbpsi

SSD & HDD price forecast

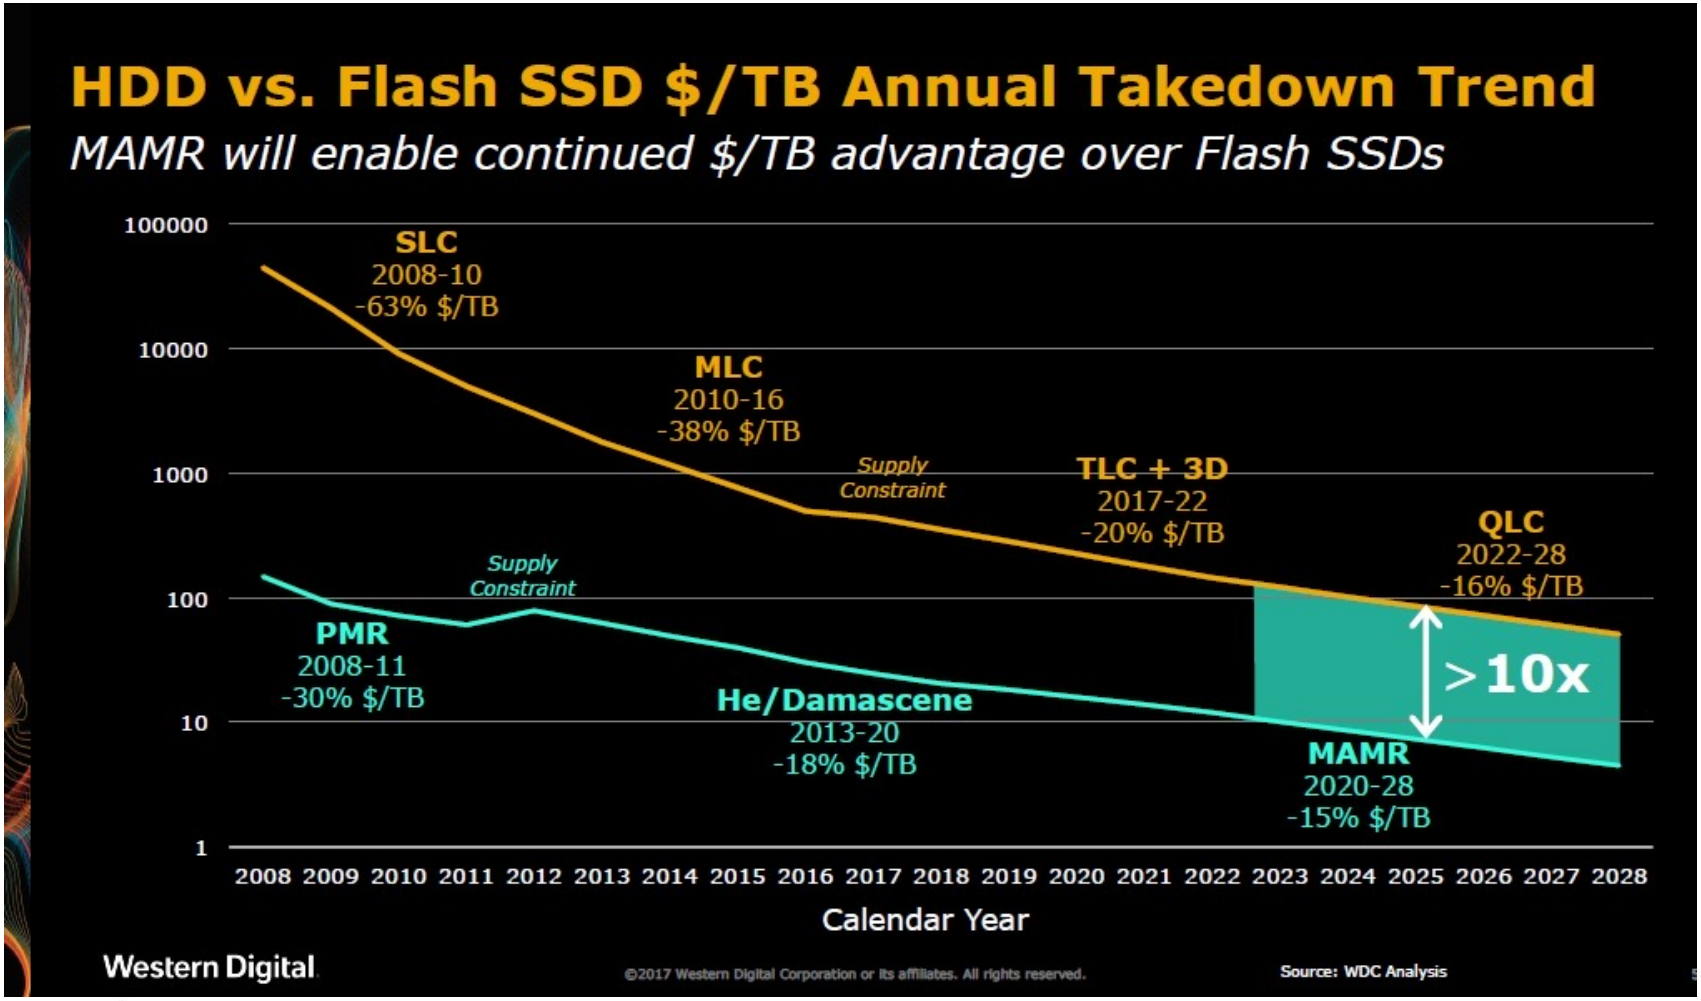

Consumer vs Enterprise

HDD

Model	Seagate Barracuda ST8000DM0004	Toshiba MG07ACA14TA	Seagate ST10000NM009G	HPE 765466-B21	HPE EG002400JWJNN
Tipus	Consumer	Enterprise	Enterprise	Enterprise	Enterprise
Capacitat (TB)	8	14	10	2	2,4
Consum (W)	6.8	7.8	9.5	7	7.1
Preu (€)	235	520	350	250	360
IOPS R/W	640	800	710	3360	5210
RPM	5400	7200	7200	10000	10000
€ / GB	0,029375	0,037142857	0,035	0,125	0,15

SSD

Model	Samsung 860 EVO	Intel Optane H10	Kingston SEDC100M	WD Gold S768T1D0D	WD Ultrastar DC SN640
Tipus	Consumer	Consumer	Enterprise	Enterprise	Enterprise
Capacitat (TB)	2	1	1,92	7,68	3,8
Consum (W)	2.2	5,8	9	12	8
Preu (€)	310	195	372	1545	750
IOPS R/W	90k / 10k	330K /250k	540K /205K	467k/ 65K	511K / 82K
Tecnologia	3D QLC NAND	3D QLC NAND	3D TLC NAND	3D TLC NAND	3D TLC NAND
€ / GB	0,155	0,195	0,19375	0,201171875	0,197368421

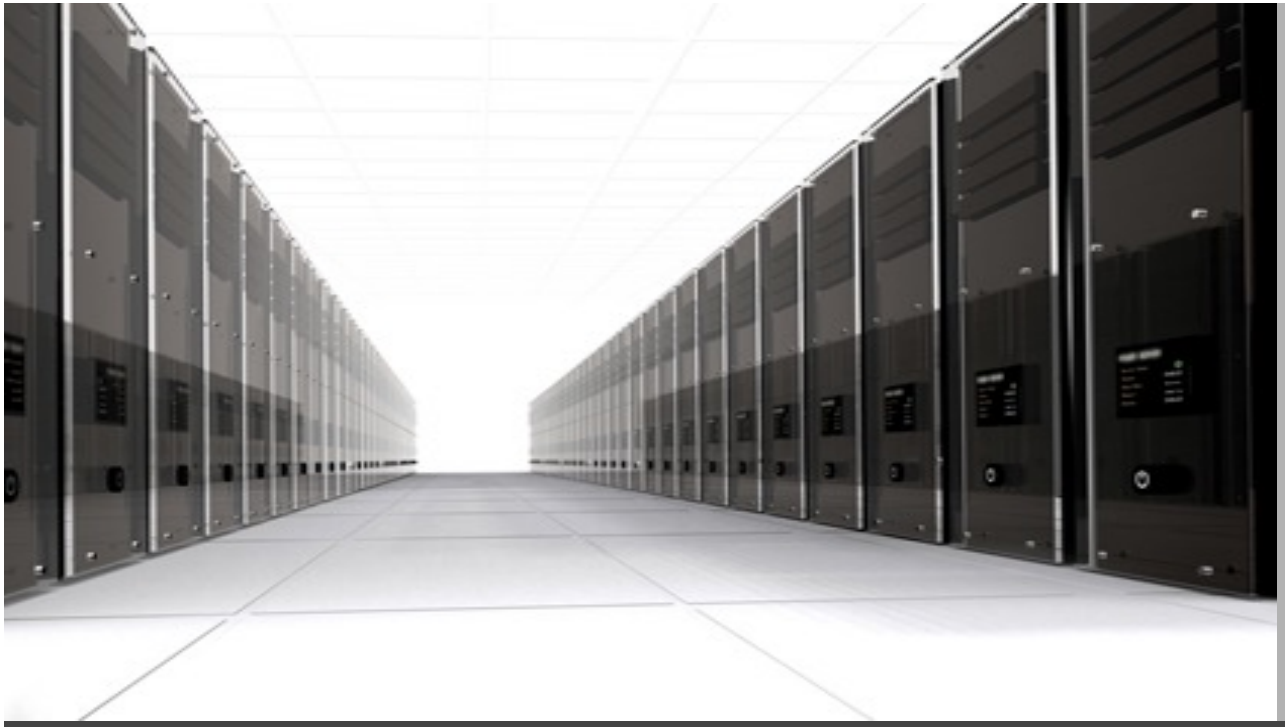

STORAGE

David López

UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH